



Achieving High Cluster Utilization with Over-Quota GPU Allocation

NetApp Solutions

Dorian Henderson, Kevin Hoke
January 14, 2021

This PDF was generated from https://docs.netapp.com/us-en/netapp-solutions/ai/osrunai_achieving_high_cluster_utilization_with_over-uota_gpu_allocation.html on May 19, 2021. Always check docs.netapp.com for the latest.

Table of Contents

Achieving High Cluster Utilization with Over-Quota GPU Allocation 1

Achieving High Cluster Utilization with Over-Quota GPU Allocation

In this section and in the sections [Basic Resource Allocation Fairness](#), and [Over-Quota Fairness](#), we have devised advanced testing scenarios to demonstrate the Run:AI orchestration capabilities for complex workload management, automatic preemptive scheduling, and over-quota GPU provisioning. We did this to achieve high cluster-resource usage and optimize enterprise-level data science team productivity in an ONTAP AI environment.

For these three sections, set the following projects and quotas:

Project	Quota
team-a	4
team-b	2
team-c	2
team-d	8

In addition, we use the following containers for these three sections:

- Jupyter Notebook: `jupyter/base-notebook`
- Run:AI quickstart: `gcr.io/run-ai-demo/quickstart`

We set the following goals for this test scenario:

- Show the simplicity of resource provisioning and how resources are abstracted from users
- Show how users can easily provision fractions of a GPU and integer number of GPUs
- Show how the system eliminates compute bottlenecks by allowing teams or users to go over their resource quota if there are free GPUs in the cluster
- Show how data pipeline bottlenecks are eliminated by using the NetApp solution when running compute-intensive jobs, such as the NetApp container
- Show how multiple types of containers are running using the system
 - Jupyter Notebook
 - Run:AI container
- Show high utilization when the cluster is full

For details on the actual command sequence executed during the testing, see [Testing Details for Section 4.8](#).

When all 13 workloads are submitted, you can see a list of container names and GPUs allocated, as shown in the following figure. We have seven training and six interactive jobs, simulating four data science teams, each with their own models running or in development. For interactive jobs, individual developers are using Jupyter Notebooks to write or debug their code. Thus, it is suitable to provision GPU fractions without using too many cluster resources.

```

root@run-deploy:~# runai list -A
NAME          STATUS  AGE  NODE  IMAGE                                TYPE      PROJECT  USER  GPUS  CREATED BY CLI  SERVICE URL(S)
b-4-gg        Running  2m   dgx1-2  gcr.io/run-ai-demo/quickstart      Train     team-b   root   2     true            http://10.61.218.134/a-1-1-jupyter,
c-5-g         Running  2m   dgx1-2  gcr.io/run-ai-demo/quickstart      Train     team-c   root   1     true            https://10.61.218.134/a-1-1-jupyter
c-4-gg        Running  2m   dgx1-1  gcr.io/run-ai-demo/quickstart      Train     team-c   root   2     true            https://10.61.218.134/a-1-1-jupyter
b-3-g         Running  2m   dgx1-1  gcr.io/run-ai-demo/quickstart      Train     team-b   root   1     true            https://10.61.218.134/a-1-1-jupyter
c-3-g02       Running  2m   dgx1-1  gcr.io/run-ai-demo/quickstart      Interactive team-c   root   0.2   true            https://10.61.218.134/a-1-1-jupyter
d-1-gggg      Running  2m   dgx1-2  gcr.io/run-ai-demo/quickstart      Train     team-d   root   4     true            https://10.61.218.134/a-1-1-jupyter
c-2-g03       Running  2m   dgx1-1  gcr.io/run-ai-demo/quickstart      Interactive team-c   root   0.3   true            https://10.61.218.134/a-1-1-jupyter
c-1-g05       Running  2m   dgx1-1  gcr.io/run-ai-demo/quickstart      Interactive team-c   root   0.5   true            https://10.61.218.134/a-1-1-jupyter
a-2-gg        Running  3m   dgx1-1  gcr.io/run-ai-demo/quickstart      Train     team-a   root   2     true            https://10.61.218.134/a-1-1-jupyter
b-2-g04       Running  3m   dgx1-2  gcr.io/run-ai-demo/quickstart      Interactive team-b   root   0.4   true            https://10.61.218.134/a-1-1-jupyter
a-1-g         Running  3m   dgx1-1  gcr.io/run-ai-demo/quickstart      Train     team-a   root   1     true            https://10.61.218.134/a-1-1-jupyter
b-1-g06       Running  3m   dgx1-2  gcr.io/run-ai-demo/quickstart      Interactive team-b   root   0.6   true            https://10.61.218.134/a-1-1-jupyter
a-1-1-jupyter Running  3m   dgx1-1  jupyter/base-notebook              Interactive team-a   root   1     true            https://10.61.218.134/a-1-1-jupyter
https://10.61.218.134/a-1-1-jupyter

```

The results of this testing scenario show the following:

- The cluster should be full: 16/16 GPUs are used.
- High cluster utilization.
- More experiments than GPUs due to fractional allocation.
- `team-d` is not using all their quota; therefore, `team-b` and `team-c` can use additional GPUs for their experiments, leading to faster time to innovation.

Next: Basic Resource Allocation Fairness

Copyright Information

Copyright © 2021 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system-without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.